

Foundations of DL

Deep Learning



ALF





Alfredo Canziani, Ritchie Ng
@alfcnz, @RitchieNg

Generative models

(Semi-) Unsupervised learning

Super resolution






Garcia (2016) srez --- github.com/david-gpu/srez

Inpainting


VAE



GAN

Yeh (2017)
Semantic Image Inpainting
with Deep Generative Models

bit.ly/DCGAN-inpainting



Caption to image

This vibrant red bird has
a pointed black beak

This bird is yellowish
orange with black wings

The bright blue bird has
a white coloured belly

Reed (2016) Generative
adversarial text to image synthesis



github.com/reedscot/icml2016



Autoencoders

Unsupervised learning / Generative models

Autoencoder



$$h = f(W_h x + b_h)$$

$$\hat{x} = g(W_x h + b_x)$$

$$x, \hat{x} \in \mathbb{R}^n$$

$$h \in \mathbb{R}^d$$

$$W_h \in \mathbb{R}^{d \times n}$$

$$W_x \in \mathbb{R}^{n \times d}$$

If "tight weights", then

$$W_x \doteq W_h^\top$$

Reconstruction losses

$$\mathcal{L} = \frac{1}{m} \sum_{j=1}^m \ell(\mathbf{x}^{(j)}, \hat{\mathbf{x}}^{(j)})$$


binary input

$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = - \sum_{i=1}^n [x_i \log(\hat{x}_i) + (1 - x_i) \log(1 - \hat{x}_i)]$$

real valued input

$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = \frac{1}{2} \|\mathbf{x} - \hat{\mathbf{x}}\|^2$$


Under-/over-complete hidden layer



\hat{x}


h

x



Denoising autoencoder

$$\tilde{x} \sim p(\tilde{x} \mid x)$$




We assume we are injecting the **same noisy distribution** we're going to observe in reality. In this way, we can learn how to robustly recover from it.

Contractive autoencoder


penalises insensitivity to
reconstruction directions

penalises sensitivity
to the **any** direction

$$\ell(\boldsymbol{x}, \hat{\boldsymbol{x}}) = \boxed{\ell_{\text{reconstruction}}} + \boxed{\lambda \|\nabla_{\boldsymbol{x}} \boldsymbol{h}\|^2}$$





Basic auto-encoder

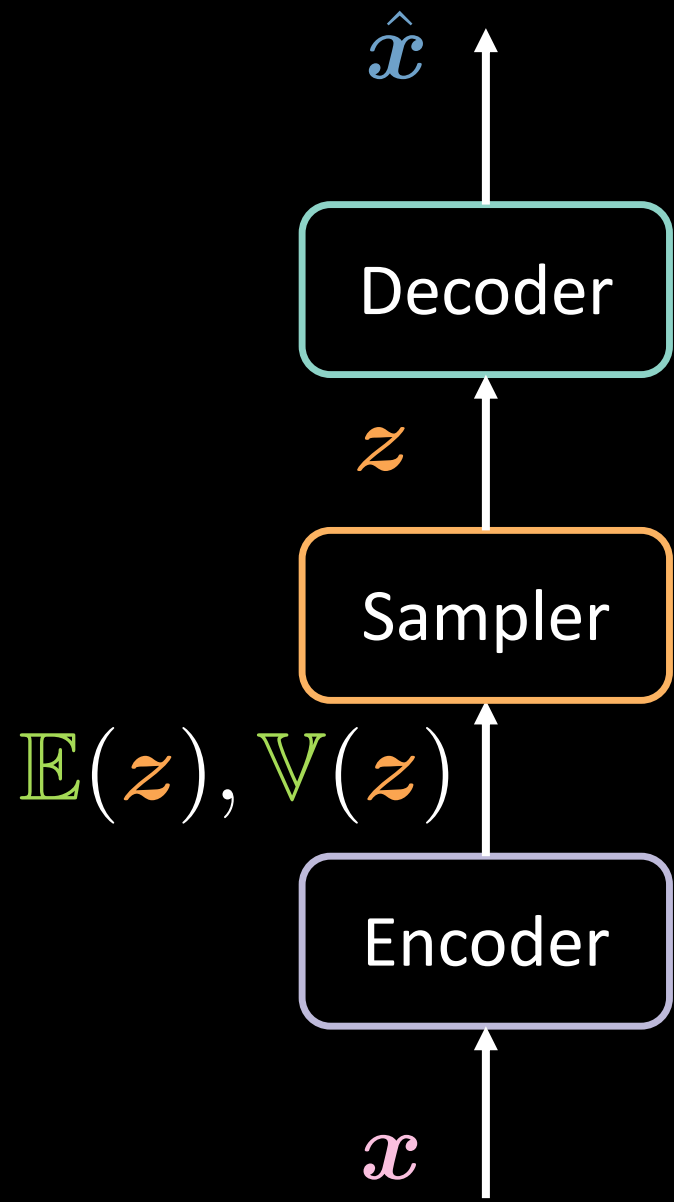


- How can we “go back”:
- learn the distribution
 - or
 - enforce some structure


Auto-encoder (recap)



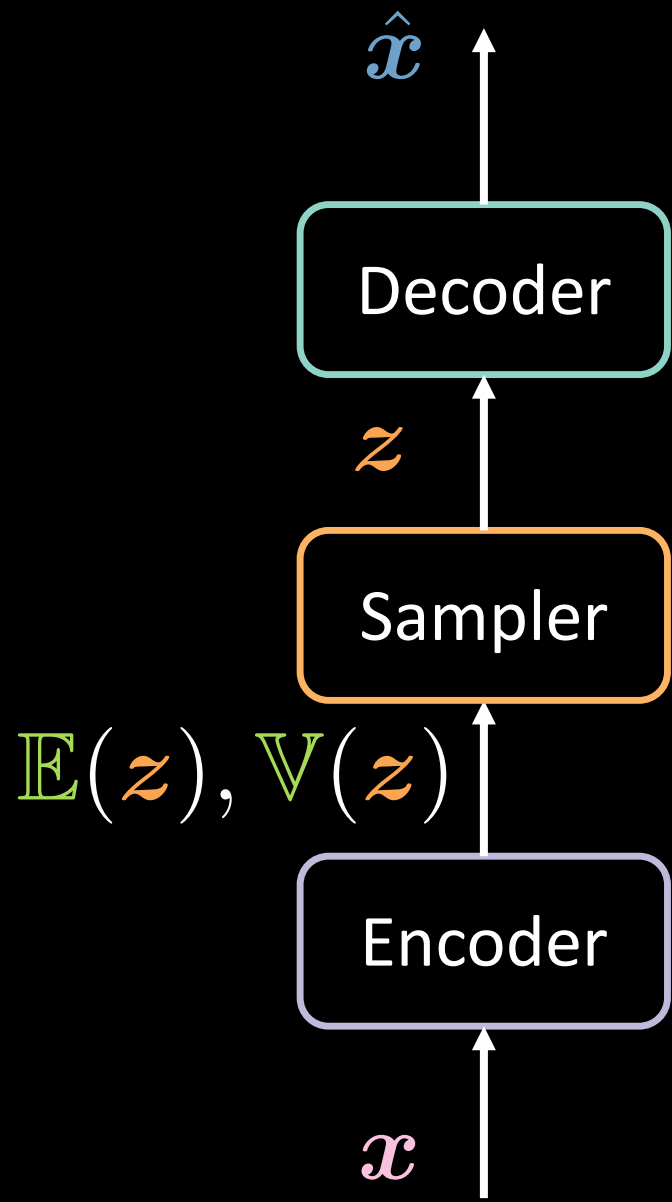
Variational auto-encoder



Classic auto-encoder




Variational auto-encoder



$$\text{decoder} : \mathcal{Z} \rightarrow \mathbb{R}^n$$
$$z \mapsto \hat{x}$$

$$\text{encoder} : \mathcal{X} \rightarrow \mathbb{R}^{2d}$$
$$x \mapsto h$$

Variational auto-encoder

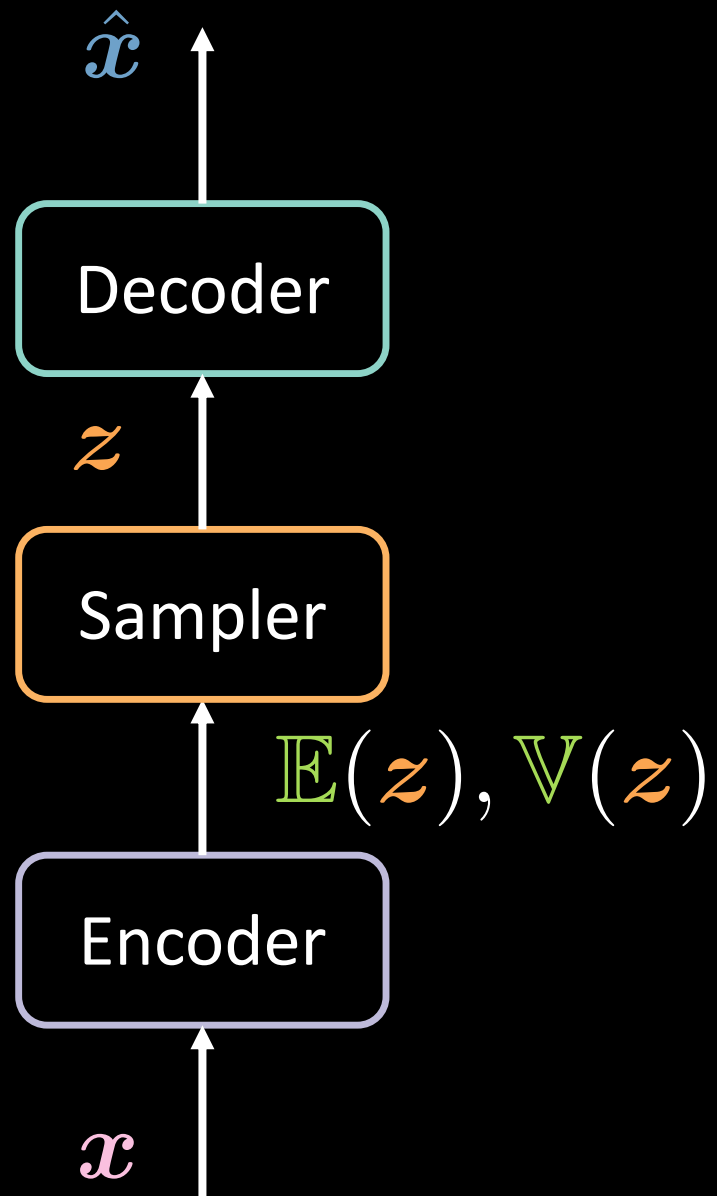


- We need to:
- learn the distribution
 - or
 - enforce some structure

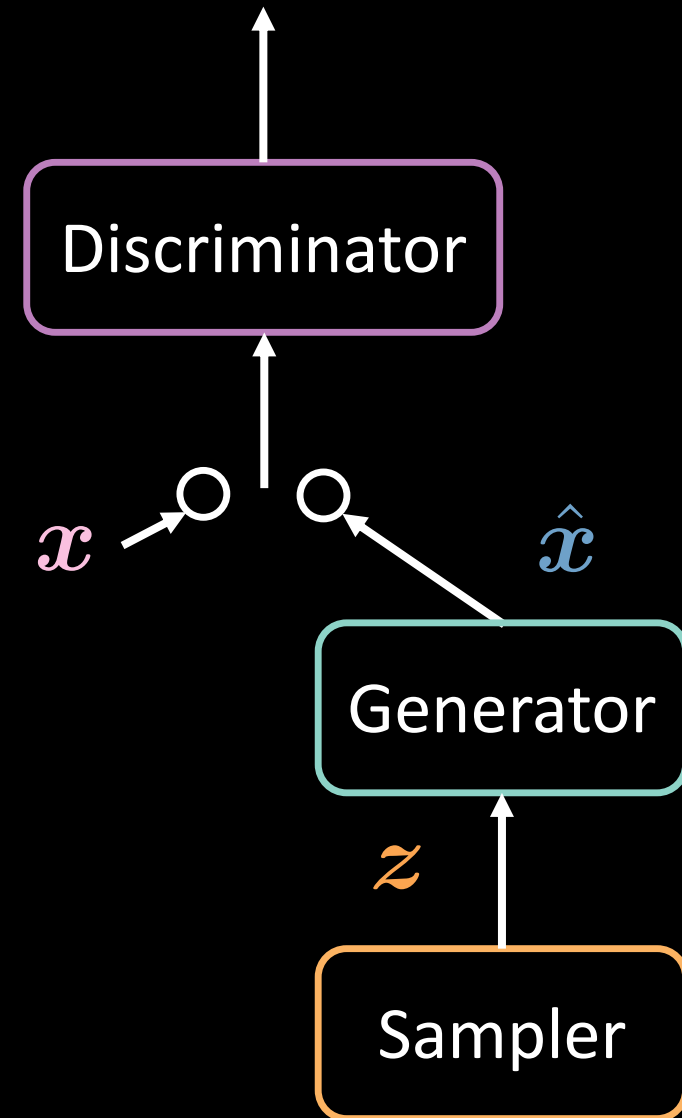
Generative adversarial nets

Unsupervised learning / Generative models

Variational auto-encoder



Generative adversarial network



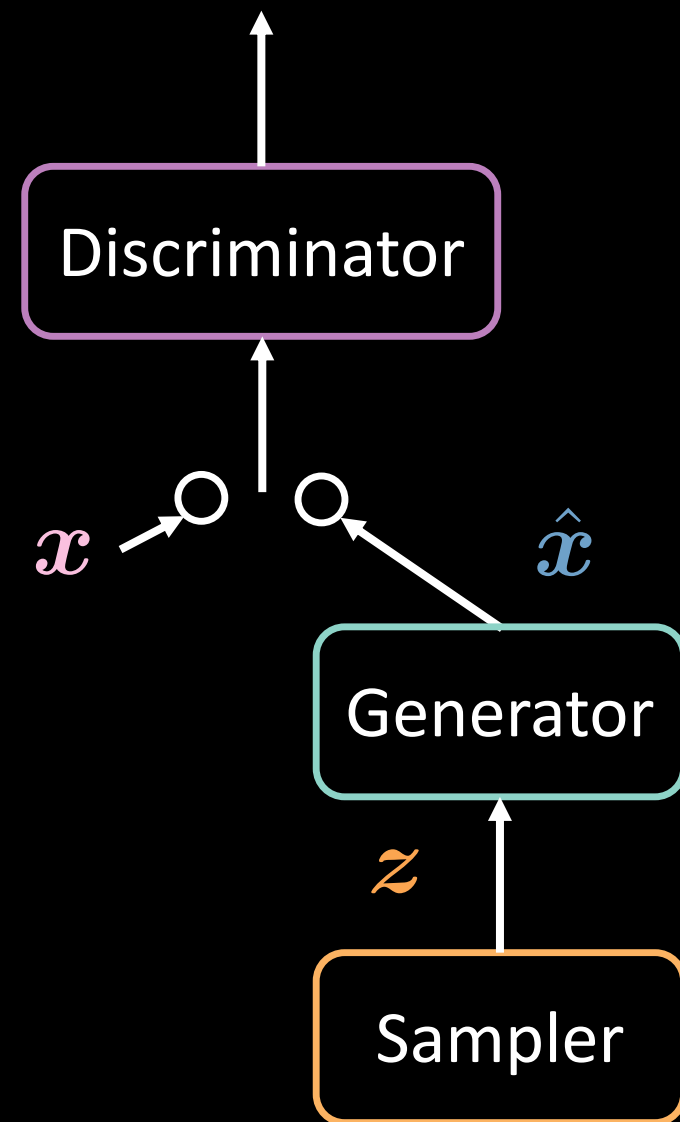
$$D : \mathbb{R}^n \rightarrow (0, 1)$$

$$x \vee \hat{x} \mapsto \ell$$


$$G : \mathcal{Z} \rightarrow \mathbb{R}^n$$

$$z \mapsto \hat{x}$$

Generative adversarial network



Generative adversarial network



Value function

$$V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \\ + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D[G(\mathbf{z})])]$$

$$\min_G \max_D V(D, G)$$